

Using Python for Data Science

Hongchang Gao


Spring 2024

Python for Data Science

- Packages:
 - Numpy
 - Scikit-learn (sklearn)
 - Pytorch
 - Tensorflow
 - Pandas
 - Matplotlib

Python for Data Science: NumPy

- NumPy: fundamental package for scientific computing
 - Large, multi-dimensional arrays and matrices
 - High-level mathematical functions



The fundamental package for scientific computing with Python

GET STARTED

NumPy v1.19.0 First Python 3 only release - Cython interface to numpy.random complete

<p>POWERFUL N-DIMENSIONAL ARRAYS</p> <p>Fast and versatile, the NumPy vectorization, indexing, and broadcasting concepts are the de-facto standards of array computing today.</p>	<p>NUMERICAL COMPUTING TOOLS</p> <p>NumPy offers comprehensive mathematical functions, random number generators, linear algebra routines, Fourier transforms, and more.</p>	<p>INTEROPERABLE</p> <p>NumPy supports a wide range of hardware and computing platforms, and plays well with distributed, GPU, and sparse array libraries.</p>
<p>PERFORMANT</p> <p>The core of NumPy is well-optimized C code. Enjoy the flexibility of Python with the speed of compiled code.</p>	<p>EASY TO USE</p> <p>NumPy's high level syntax makes it accessible and productive for programmers from any background or experience level.</p>	<p>OPEN SOURCE</p> <p>Distributed under a liberal BSD license, NumPy is developed and maintained publicly on GitHub by a vibrant, responsive, and diverse community.</p>

Python for Data Science: NumPy

- NumPy: fundamental package for scientific computing
 - 1-D array: `arr = np.array([1, 2, 3, 4, 5])`
 - 2-D array `arr = np.array([[1, 2, 3], [4, 5, 6]])`
 - Array indexing: `arr[2] + arr[3]`
 - Mathematic operation

```
>>> a = np.array([1,2,3])
>>> b = np.array([0,1,0])
>>> np.inner(a, b)
2
```

Python for Data Science: scikit-learn (sklearn)


- Machine learning package in python

Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more...




Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more...



Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more...



Python for Data Science: scikit-learn

- It includes various datasets for machine learning tasks

<code>datasets.fetch_20newsgroups(*[, data_home, ...])</code>	Load the filenames and data from the 20 newsgroups dataset (classification).
<code>datasets.fetch_20newsgroups_vectorized(*[, ...])</code>	Load and vectorize the 20 newsgroups dataset (classification).
<code>datasets.fetch_california_housing(*[, ...])</code>	Load the California housing dataset (regression).
<code>datasets.fetch_covtype(*[, data_home, ...])</code>	Load the covtype dataset (classification).
<code>datasets.fetch_kddcup99(*[, subset, ...])</code>	Load the kddcup99 dataset (classification).
<code>datasets.fetch_lfw_pairs(*[, subset, ...])</code>	Load the Labeled Faces in the Wild (LFW) pairs dataset (classification).
<code>datasets.fetch_lfw_people(*[, data_home, ...])</code>	Load the Labeled Faces in the Wild (LFW) people dataset (classification).
<code>datasets.fetch_olivetti_faces(*[, ...])</code>	Load the Olivetti faces data-set from AT&T (classification).
<code>datasets.fetch_openml([name, version, ...])</code>	Fetch dataset from openml by name or dataset id.
<code>datasets.fetch_rcv1(*[, data_home, subset, ...])</code>	Load the RCV1 multilabel dataset (classification).
<code>datasets.fetch_species_distributions(*[, ...])</code>	Loader for species distribution dataset from Phillips et.

Python for Data Science: scikit-learn

- It includes various state-of-the-art machine learning models

<code>cluster.AffinityPropagation(*[, damping, ...])</code>	Perform Affinity Propagation Clustering of data.
<code>cluster.AgglomerativeClustering([...])</code>	Agglomerative Clustering
<code>cluster.Birch(*[, threshold, ...])</code>	Implements the Birch clustering algorithm.
<code>cluster.DBSCAN([eps, min_samples, metric, ...])</code>	Perform DBSCAN clustering from vector array or distance matrix.
<code>cluster.FeatureAgglomeration([n_clusters, ...])</code>	Agglomerate features.
<code>cluster.KMeans([n_clusters, init, n_init, ...])</code>	K-Means clustering.
<code>cluster.MinibatchKMeans([n_clusters, init, ...])</code>	Mini-Batch K-Means clustering.
<code>cluster.MeanShift(*[, bandwidth, seeds, ...])</code>	Mean shift clustering using a flat kernel.
<code>cluster.OPTICS(*[, min_samples, max_eps, ...])</code>	Estimate clustering structure from vector array.
<code>cluster.SpectralClustering([n_clusters, ...])</code>	Apply clustering to a projection of the normalized Laplacian.
<code>cluster.SpectralBiclustering([n_clusters, ...])</code>	Spectral biclustering (Kluger, 2003).
<code>cluster.SpectralCoclustering([n_clusters, ...])</code>	Spectral Co-Clustering algorithm (Dhillon, 2001).

Unsupervised

Python for Data Science: scikit-learn

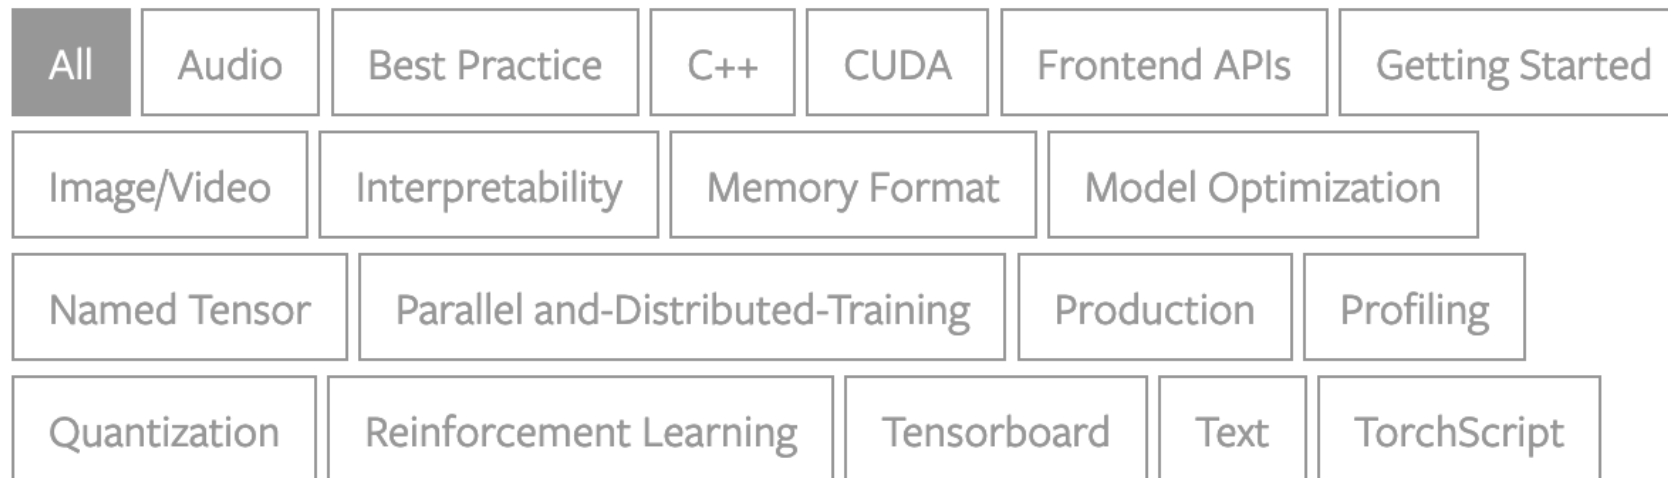
- It includes various state-of-the-art machine learning models

<code>linear_model.ElasticNet([alpha, l1_ratio, ...])</code>	Linear regression with combined L1 and L2 priors as regularizer.
<code>linear_model.ElasticNetCV(*[, l1_ratio, ...])</code>	Elastic Net model with iterative fitting along a regularization path.
<code>linear_model.Lars(*[, fit_intercept, ...])</code>	Least Angle Regression model a.k.a.
<code>linear_model.LarsCV(*[, fit_intercept, ...])</code>	Cross-validated Least Angle Regression model.
<code>linear_model.Lasso([alpha, fit_intercept, ...])</code>	Linear Model trained with L1 prior as regularizer (aka the Lasso)
<code>linear_model.LassoCV(*[, eps, n_alphas, ...])</code>	Lasso linear model with iterative fitting along a regularization path.
<code>linear_model.LassoLars([alpha, ...])</code>	Lasso model fit with Least Angle Regression a.k.a.
<code>linear_model.LassoLarsCV(*[, fit_intercept, ...])</code>	Cross-validated Lasso, using the LARS algorithm.
<code>linear_model.LassoLarsIC([criterion, ...])</code>	Lasso model fit with Lars using BIC or AIC for model selection
<code>linear_model.OrthogonalMatchingPursuit(*[, ...])</code>	Orthogonal Matching Pursuit model (OMP).
<code>linear_model.OrthogonalMatchingPursuitCV(*)</code>	Cross-validated Orthogonal Matching Pursuit model (OMP).

Supervised

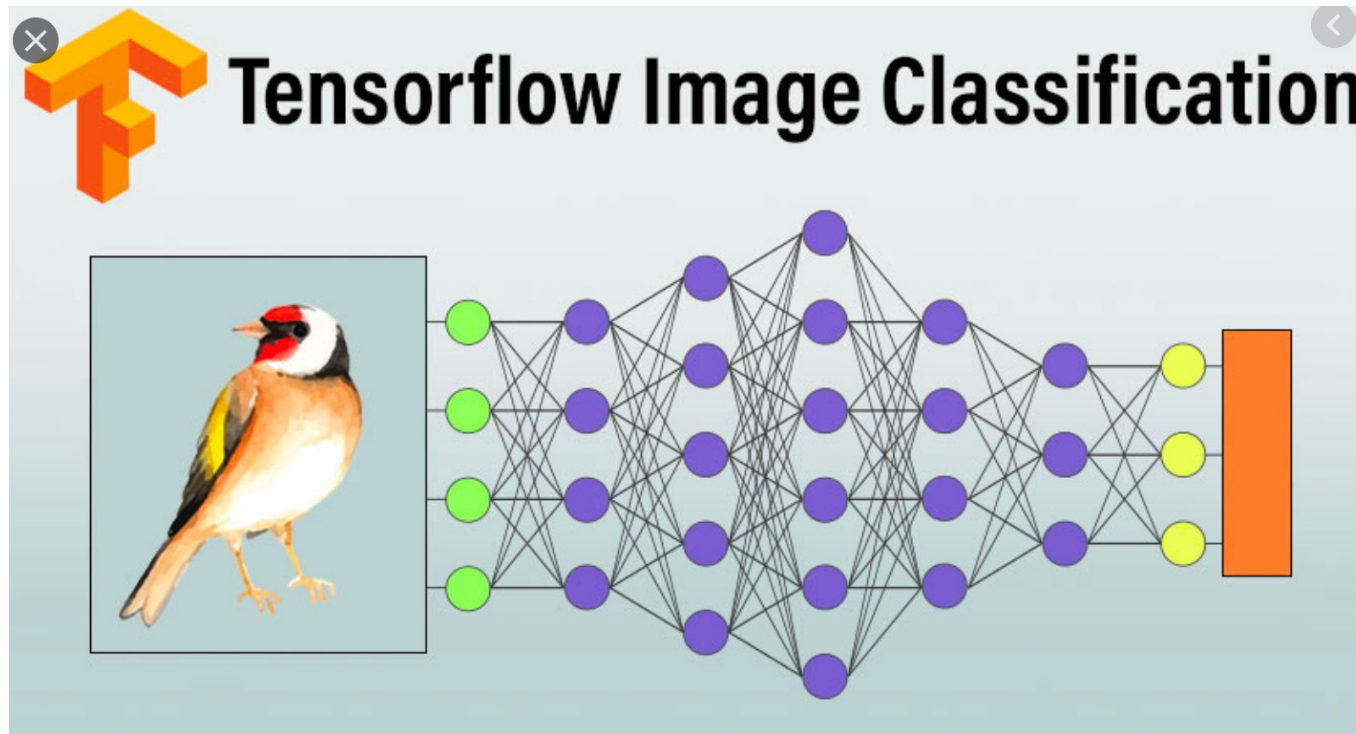
Python for Data Science: PyTorch

- Deep Learning tools in Python, developed by Facebook



Python for Data Science: Tensorflow

- Deep Learning tools in Python, developed by Google



Python for Data Science: Pandas

- A Python package for working with datasets
- Has functions for
 - Analyze data
 - Clean data
 - Explore data
 - Manipulate data


```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 169 entries, 0 to 168  
Data columns (total 4 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Duration    169 non-null    int64  
1   Pulse       169 non-null    int64  
2   Maxpulse    169 non-null    int64  
3   Calories    164 non-null    float64  
dtypes: float64(1), int64(3)  
memory usage: 5.4 KB  
None
```




```
import pandas as pd
```

```
df = pd.read_json('data.json')
```



Python for Data Science: Matplotlib

- A Python package for visualizing data



Pandas

- A python library for working with data sets
- Has functions for
 - Analyzing data
 - Cleaning data
 - Exploring data
 - Manipulating data

Pandas

- DataFrame

- A 2-dimensional data structure, like a 2-dimensional array, or a table with rows and columns.

```
import pandas as pd
```

```
data = {  
    "calories": [420, 380, 390],  
    "duration": [50, 40, 45]  
}
```

```
df = pd.DataFrame(data)
```

```
print(df)  
print(type(df))
```

	calories	duration
0	420	50
1	380	40
2	390	45

<class 'pandas.core.frame.DataFrame'>

Pandas

- Load CSV file

	A	B	C	D	E	F	G	H	I
1	Car	MPG	Cylinders	Displacemen	Horsepower	Weight	Acceleration	Model	Origin
2	Chevrolet Chevelle	18	8	307	130	3504	12	70	US
3	Buick Skylark	15	8	350	165	3693	11.5	70	US
4	Plymouth Satellite	18	8	318	150	3436	11	70	US
5	AMC Rebel Sport	16	8	304	150	3433	12	70	US
6	Ford Torino	17	8	302	140	3449	10.5	70	US
7	Ford Galaxie	15	8	429	198	4341	10	70	US
8	Chevrolet Impala	14	8	454	220	4354	9	70	US
9	Plymouth Fury	14	8	440	215	4312	8.5	70	US
10	Pontiac Catalina	14	8	455	225	4425	10	70	US
11	AMC Ambassador	15	8	390	190	3850	8.5	70	US
12	Citroen DS-2	0	4	133	115	3090	17.5	70	Europe
13	Chevrolet Chevelle	0	8	350	165	4142	11.5	70	US
14	Ford Torino (LTD)	0	8	351	153	4034	11	70	US
15	Plymouth Satellite	0	8	383	175	4166	10.5	70	US
16	AMC Rebel Sport	0	8	360	175	3850	11	70	US
17	Dodge Challenger	15	8	383	170	3563	10	70	US

```
import pandas as pd

df = pd.read_csv("cars.csv")

print(df)
```

	Car	MPG	Cylinders	Displacement	\
0	Chevrolet Chevelle Malibu	18.0	8	307.0	
1	Buick Skylark 320	15.0	8	350.0	
2	Plymouth Satellite	18.0	8	318.0	
3	AMC Rebel SST	16.0	8	304.0	
4	Ford Torino	17.0	8	302.0	
5	Ford Galaxie 500	15.0	8	429.0	
6	Chevrolet Impala	14.0	8	454.0	
7	Plymouth Fury iii	14.0	8	440.0	
8	Pontiac Catalina	14.0	8	455.0	
9	AMC Ambassador DPL	15.0	8	390.0	
10	Citroen DS-21 Pallas	0.0	4	133.0	
11	Chevrolet Chevelle Concours (sw)	0.0	8	350.0	
12	Ford Torino (sw)	0.0	8	351.0	

Pandas

- Get rows

```
print("===The first row===")  
  
print(df.iloc[0])
```

```
===The first row===  
Car          Chevrolet Chevelle Malibu  
MPG          18  
Cylinders    8  
Displacement 307  
Horsepower   130  
Weight       3504  
Acceleration 12  
Model        70  
Origin       US  
Name: 0, dtype: object
```

```
print("===The last row===")  
  
print(df.iloc[-1])
```

```
===The first row===  
Car          Chevy S-10  
MPG          31  
Cylinders    4  
Displacement 119  
Horsepower   82  
Weight       2720  
Acceleration 19.4  
Model        82  
Origin       US  
Name: 405, dtype: object
```

Pandas

- Get columns

```
print("===The first column===")  
  
print(df.iloc[:, 0])
```

```
===The first column===  
0      Chevrolet Chevelle Malibu  
1      Buick Skylark 320  
2      Plymouth Satellite  
3      AMC Rebel SST  
4      Ford Torino  
5      Ford Galaxie 500  
6      Chevrolet Impala  
7      Plymouth Fury iii  
8      Pontiac Catalina  
9      AMC Ambassador DPL  
10     Citroen DS-21 Pallas
```

```
print("===The last column===")  
  
print(df.iloc[:, -1])
```

```
===The last column===  
0      US  
1      US  
2      US  
3      US  
4      US  
5      US  
6      US  
7      US  
8      US  
9      US  
10     Europe
```

Pandas

- Get multiple rows or columns

```
print(df.iloc[1, [2,3]]) # row 1; columns 2, 3
```

```
Cylinders      8  
Displacement  350  
Name: 1, dtype: object
```

```
print(df.iloc[[1,2], 3]) # rows 1, 2; column 3
```

```
1    350.0  
2    318.0  
Name: Displacement, dtype: float64
```

Pandas

- Convert DataFrame to numpy Array

```
print(df.to_numpy())
```

```
[['Chevrolet Chevelle Malibu' 18.0 8 ... 12.0 70 'US']  
 ['Buick Skylark 320' 15.0 8 ... 11.5 70 'US']  
 ['Plymouth Satellite' 18.0 8 ... 11.0 70 'US']  
 ...  
 ['Dodge Rampage' 32.0 4 ... 11.6 82 'US']  
 ['Ford Ranger' 28.0 4 ... 18.6 82 'US']  
 ['Chevy S-10' 31.0 4 ... 19.4 82 'US']]
```

```
print(df.iloc[:, 1:-1].to_numpy())
```

```
[ [ 18.      8.    307.    ... 3504.     12.     70. ]  
  [ 15.      8.    350.    ... 3693.     11.5    70. ]  
  [ 18.      8.    318.    ... 3436.     11.     70. ]  
  ...  
  [ 32.      4.    135.    ... 2295.     11.6    82. ]  
  [ 28.      4.    120.    ... 2625.     18.6    82. ]  
  [ 31.      4.    119.    ... 2720.     19.4    82. ]]
```


Matplotlib

- A Python package for visualization
 - Plot: draws a line from point to point

```
import matplotlib.pyplot as plt
import numpy as np

x_coord = np.array([1, 8])
y_coord = np.array([3, 10])

plt.plot(x_coord, y_coord)
plt.show()
```




Matplotlib

- Scatter
 - plot one dot for each point

```
import matplotlib.pyplot as plt
import numpy as np

x = np.array([5, 7, 8, 7, 2, 17, 2, 9, 4, 11, 12, 9, 6])
y = np.array([99, 86, 87, 88, 111, 86, 103, 87, 94, 78, 77, 85, 86])

plt.scatter(x, y)
plt.show()
```




Matplotlib

- Bar plot
 - draw bar graph

```
import matplotlib.pyplot as plt
import numpy as np

x = np.array(["A", "B", "C", "D"])
y = np.array([3, 8, 1, 10])

plt.bar(x,y)
plt.show()
```




Matplotlib

- Histogram plot
 - the number of observations within each given interval

```
import matplotlib.pyplot as plt
import numpy as np

x = np.random.normal(170, 10, 250)

plt.hist(x)
plt.show()
```



NumPy

- Create a NumPy nd-array Object
 - pass a list, tuple or any array-like object into the array() method

```
import numpy as np

a = np.array(42)
b = np.array([1, 2, 3, 4, 5])
c = np.array([[1, 2, 3], [4, 5, 6]])
d = np.array([[[1, 2, 3], [4, 5, 6]], [[1, 2, 3], [4, 5, 6]]])

print(a.ndim)    0
print(b.ndim)    1
print(c.ndim)    2
print(d.ndim)    3
```

NumPy

- Access array elements

```
import numpy as np

arr = np.array([1, 2, 3, 4])

print(arr[0])
print(arr[-1])
print(arr[0:2])
```

```
1
4
[1 2]
```

```
import numpy as np

arr = np.array([[1,2,3], [6,7,8]])

print('1st row: ', arr[0, :])
print('1st col: ', arr[:, 0])
print('2nd element on 1st row: ', arr[0, 1])
```

```
1st row:  [1 2 3]
1st col:  [1 6]
2nd element on 1st row:  2
```

NumPy

- Array slicing:
 - taking elements from one given index to another given index

```
import numpy as np

arr = np.array([1, 2, 3, 4, 5, 6, 7])

print(arr[1:5])
print(arr[1:])
print(arr[:5])
print(arr[:])
print(arr[1:-2])
print(arr[1:5:2])
```

```
[2 3 4 5]
[2 3 4 5 6 7]
[1 2 3 4 5]
[1 2 3 4 5 6 7]
[2 3 4 5]
[2 4]
```

NumPy

- Joining NumPy Arrays

- putting contents of two or more arrays in a single array

```
import numpy as np

arr1 = np.array([[1, 2], [3, 4]])
arr2 = np.array([[5, 6], [7, 8]])
print(arr1)
print(arr2)

new_arr = np.concatenate((arr1, arr2), axis=1)
print(new_arr)

new_arr = np.concatenate((arr1, arr2), axis=0)
print(new_arr)
```

```
[[1 2]
 [3 4]
 [5 6]
 [7 8]]
[[1 2 5 6]
 [3 4 7 8]]
[[1 2]
 [3 4]
 [5 6]
 [7 8]]
```