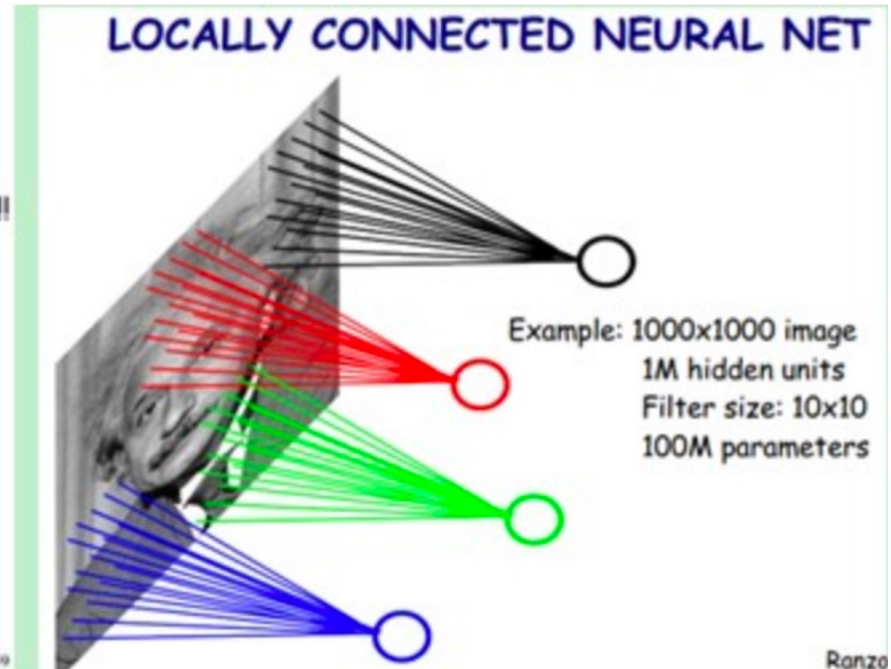
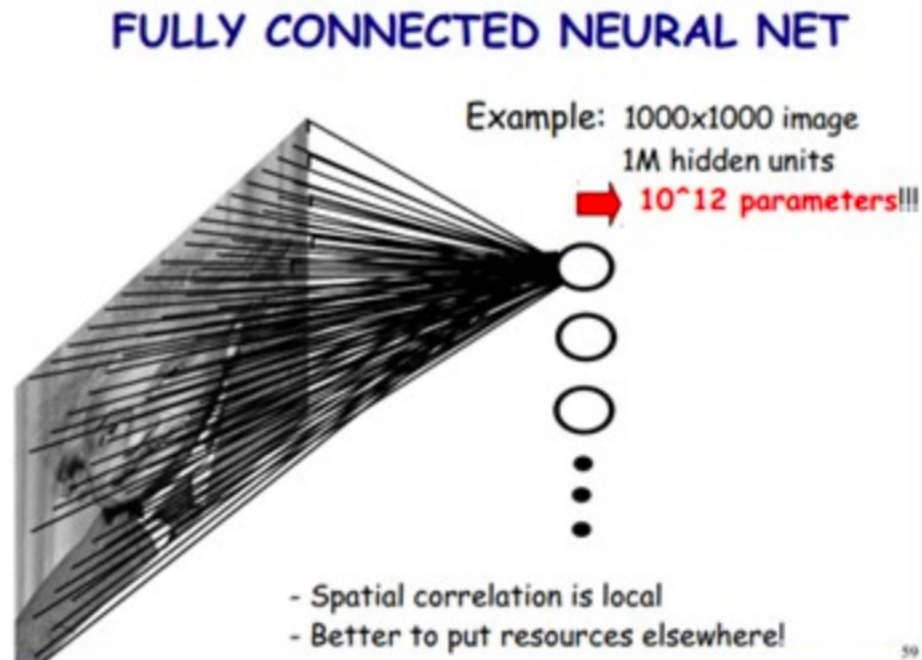


Convolutional Neural Network

Spring 2024


Hongchang Gao

Background



Background

- **Feature extractor:** Learn useful features from images for prediction
 - 1. Sobel Filter - Weights to Detect Horizontal Edges



*


1	2	1
0	0	0
-1	-2	-1



Horizontal Edge
(absolute value)⁸



Background

- **Feature extractor:** Learn useful features from images for prediction
 - 1. Sobel Filter - Weights to Detect Vertical Edges



*

1	0	-1
2	0	-2
1	0	-1



Vertical Edge₇
(absolute value)

Convolution

- Matrix Inner Product


- $\mathbf{A} = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$ and $\mathbf{B} = \begin{bmatrix} 5 & 7 \\ 6 & 8 \end{bmatrix}$.

- Inner product:

$$\langle \mathbf{A}, \mathbf{B} \rangle = \sum_i \sum_j a_{ij} b_{ij} = 70.$$

- Property: $\langle \mathbf{A}, \mathbf{B} \rangle = \langle \text{vec}(\mathbf{A}), \text{vec}(\mathbf{B}) \rangle$.

Convolution



Convolution

Input Image
5×5

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Filter (Kernel)
3×3

1	0	1
0	1	0
1	0	1

Convolution:

4	3	4
2	4	3
2	3	4

Result
3×3

The value **4** is the inner product of the patch

1	1	1
0	1	1
0	0	1

and the filter

1	0	1
0	1	0
1	0	1

Convolution

Input Image
5×5

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Filter (Kernel)
3×3

1	0	1
0	1	0
1	0	1

Convolution:

4	3	4
2	4	3
2	3	4

Result
3×3

The value **3** is the inner product of the patch

1	1	0
1	1	1
0	1	1

and the filter

1	0	1
0	1	0
1	0	1

Convolution

Input Image
5×5

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Filter (Kernel)
3×3

1	0	1
0	1	0
1	0	1

Convolution:

4	3	4
2	4	3
2	3	4

Result
3×3

The value **4** is the inner product of the patch

1	1	1
1	1	0
1	0	0

and the filter

1	0	1
0	1	0
1	0	1

Convolution

- Question: How many 3x3 patches in following two images?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0


1	1	1	0
0	1	1	1
0	1	1	0
1	1	0	0

Dimensions:

- Input: $d_1 \times d_2$
- Filter: $k_1 \times k_2$
- Output: $(d_1 - k_1 + 1) \times (d_2 - k_2 + 1)$

Zero Padding


- Problem: the output is smaller than the input



- Zero padding
 - Keep the next layer's width and height consistent with the previous
 - Keep the information around the border of the image



Zero Padding

- Add a “boarder” of all-zeros.
- Increase the input shape:
 - From $d_1 \times d_2$ to $(d_1 + 2) \times (d_2 + 2)$.
- If the filter is 3×3 , the output is $d_1 \times d_2$



Stride

- Stride
 - The filter moves K step each time, $K \geq 1$



Stride

- Stride
 - The filter moves K step each time, $K \geq 1$

Dimensions:

- Input: $d_1 \times d_2$
- Filter: $k_1 \times k_2$
- Stride: s
- Output: $\left(\left\lfloor \frac{d_1 - k_1}{s} \right\rfloor + 1 \right) \times \left(\left\lfloor \frac{d_2 - k_2}{s} \right\rfloor + 1 \right)$

Summary

- Convolution:
 - Feature map, filter/kernel
 - Zero padding
 - Stride

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Feature map

1	0	1
0	1	0
1	0	1


Filter/kernel

4	3	4
2	4	3
2	3	4

Feature map

Summary


- Traditional **Feature extractor**:
 - hand-crafted filters
 - a lot of expertise
 - Limited filters




$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$




$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



Summary

- Traditional **Feature extractor**:
 - hand-crafted filters
 - a lot of expertise
 - Limited filters

Can we learn filters automatically?




Convolutional Neural Network

- Convolutional layer



$$x_j^\ell = f\left(\sum_i \text{conv}(x_i^{\ell-1}, k_{ij}^\ell) + b_j^\ell\right)$$

x_j^ℓ is the j -th feature map in the ℓ -th layer,



k_{ij}^ℓ is the convolutional kernel in the ℓ -th layer



Feature map/channel



filters



Feature map/channel

3 input feature maps
1 output feature maps
3x1 filters

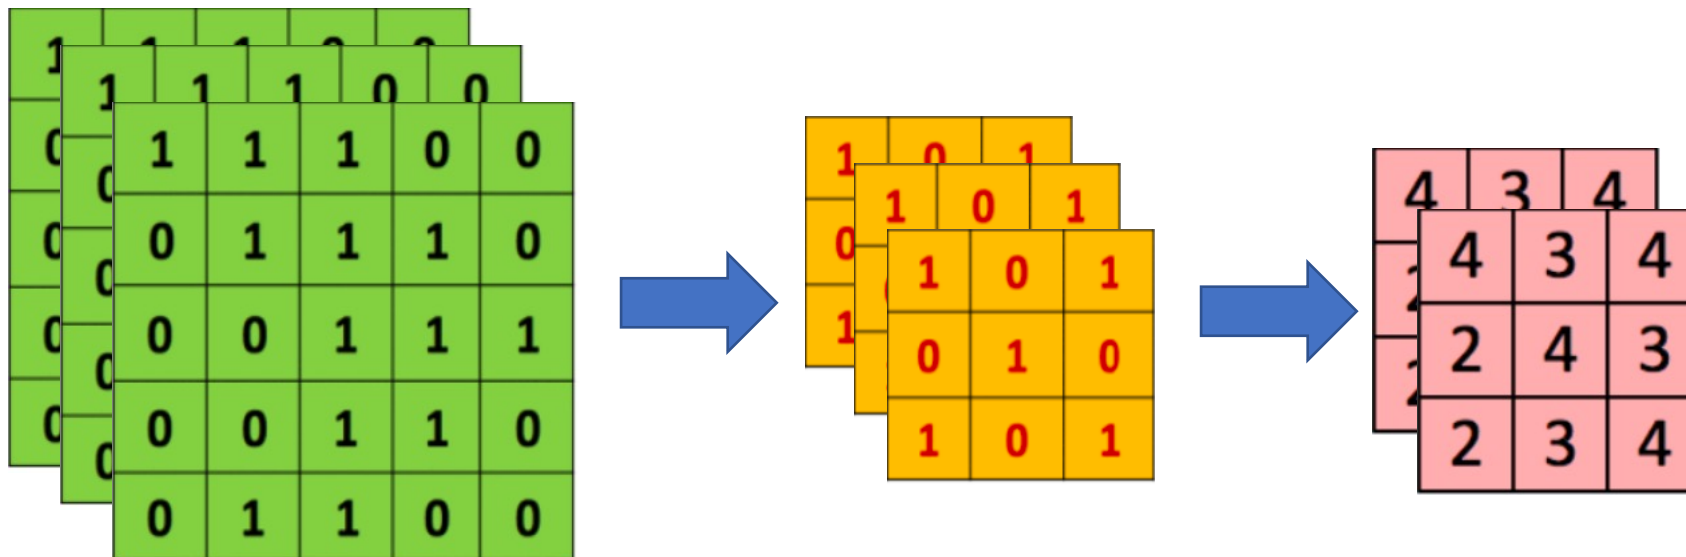
Convolutional Neural Network

- Convolutional layer

$$x_j^\ell = f\left(\sum_i \text{conv}(x_i^{\ell-1}, k_{ij}^\ell) + b_j^\ell\right)$$

x_j^ℓ is the j -th feature map in the ℓ -th layer,

k_{ij}^ℓ is the convolutional kernel in the ℓ -th layer



3 input feature maps
2 output feature maps
3x2 filters

Convolutional Neural Network

- Question:

Input features: $5 \times 32 \times 32$


Convolution kernel: $8 \times 5 \times 3 \times 3$

$$x_j^\ell = f\left(\sum_i conv(x_i^{\ell-1}, k_{ij}^\ell) + b_j^\ell\right)$$

- ▶ How many input channels are there?
- ▶ How many output channels are there?
- ▶ How many trainable weights are there?

Convolutional Neural Network


- Pooling layer
 - Reduce the dimensionality of each feature map
 - Max pooling, mean pooling



- Pool size = 2×2 .
- No overlap (stride = 2×2).

Convolutional Neural Network

- Pooling layer
 - Reduce the dimensionality of each feature map
 - Max pooling, mean pooling



- Pool size = 2×2 .
- No overlap (stride = 2×2).

Summary


- CNN
 - Convolutional layer

$$x_j^\ell = f\left(\sum_i conv(x_i^{\ell-1}, k_{ij}^\ell) + b_j^\ell\right)$$

- Pooling layer
 - Reduce dimensionality


Example

- Convolution + pooling 1




Example

- Convolution + pooling 2



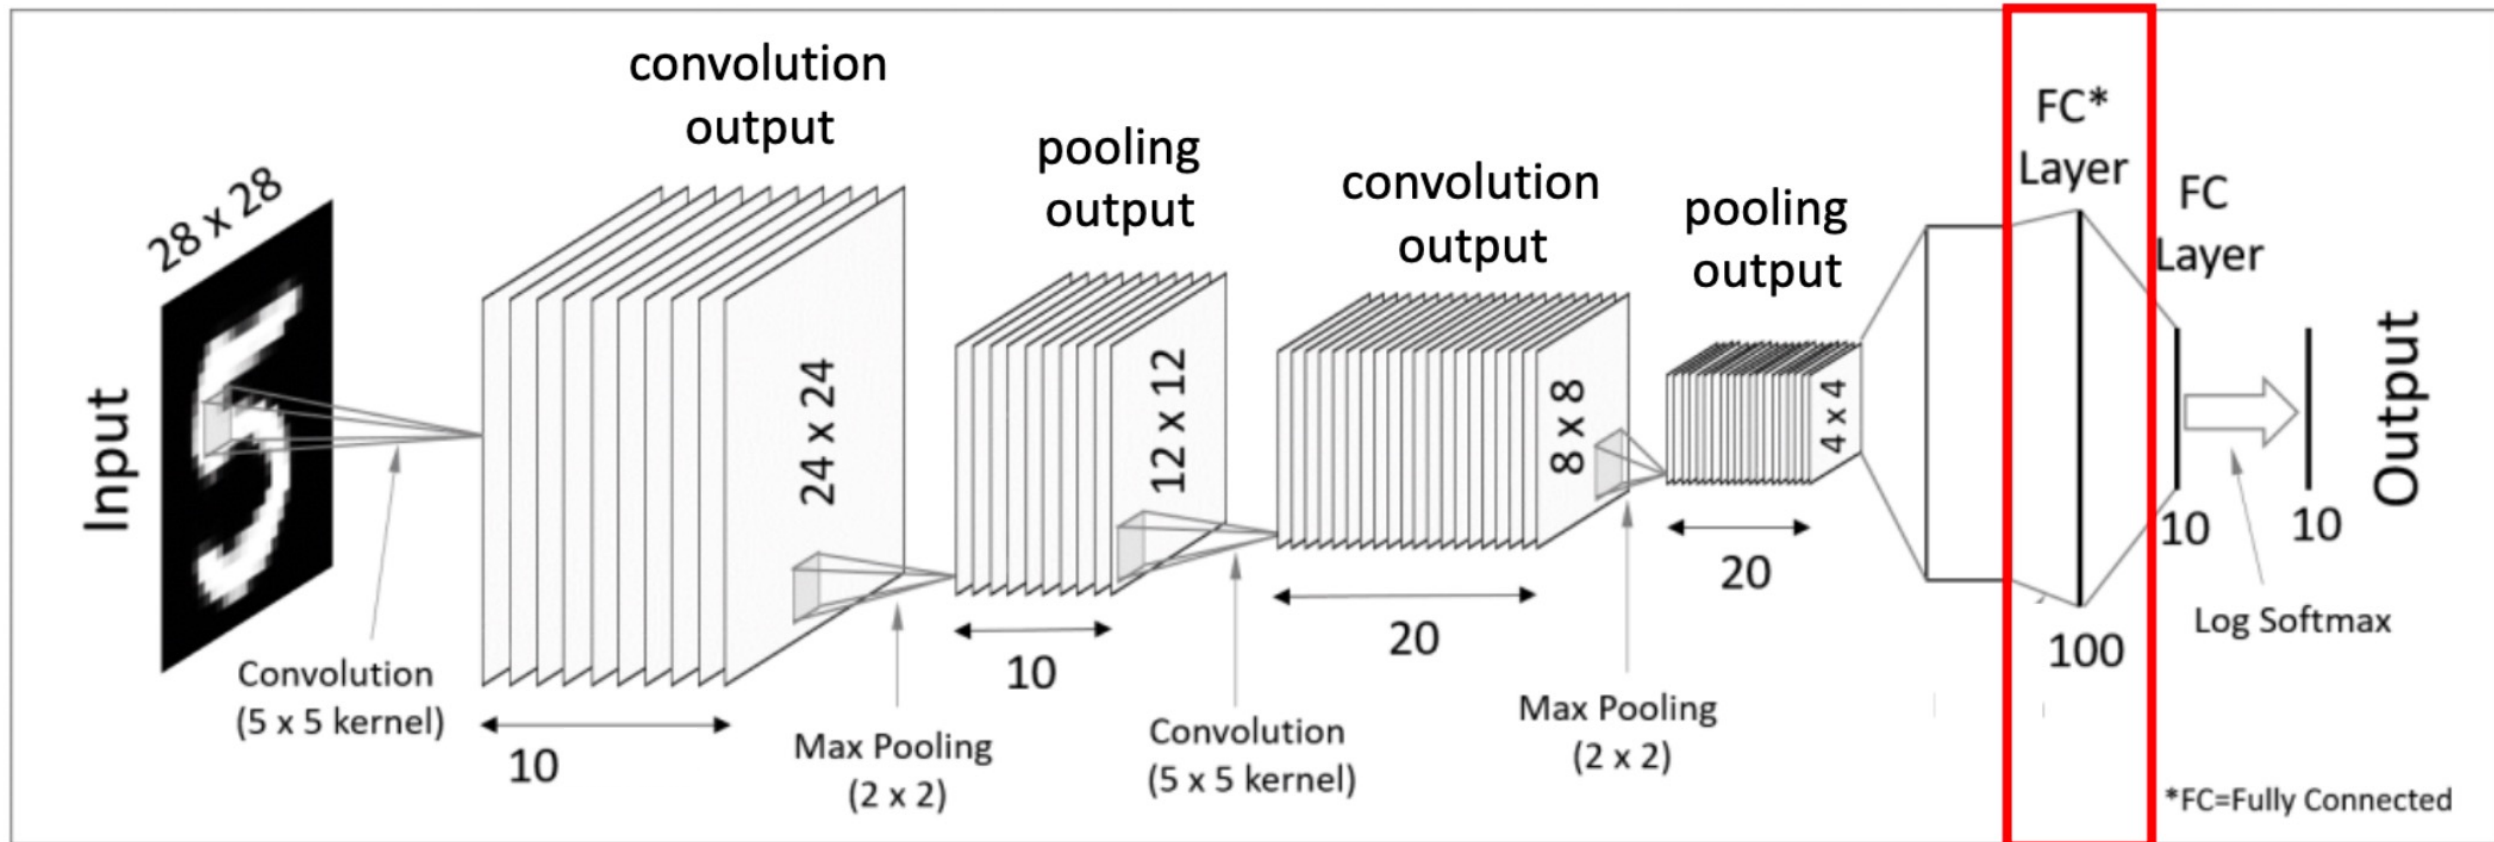
Example

- Flatten the 20 feature maps (4x4)



Example

- Add one fully-connected layer



Question

- How many model parameters?

